Highlights

- A cost-constrained attribute reduction framework for label distribution learning.
- Non-monotone granular-ball dependency guides minimum-cost subset search.
- An internal feasibility guard certifies zero-cost answers on degenerate domains.
- Acquisition costs drop by two orders of magnitude without accuracy loss.
- A ridge predictor revives nine benchmarks; low-cost subsets dominate full sets.